

# Optimizing Three-Dimensional Stencil-Operations on Heterogeneous Computing Environments

Nina Herrmann<sup>1\*†</sup>, Justus Dieckmann<sup>1†</sup> and Herbert Kuchen<sup>1</sup>

<sup>1\*</sup>Practical Computer Science, University of Münster, Leonardo-Campus  
3, Münster, 48149, Germany.

\*Corresponding author(s). E-mail(s): [nina.herrmann@uni-muenster.de](mailto:nina.herrmann@uni-muenster.de);

Contributing authors: [justus.dieckmann@uni-muenster.de](mailto:justus.dieckmann@uni-muenster.de);

[kuchen@uni-muenster.de](mailto:kuchen@uni-muenster.de);

<sup>†</sup>These authors contributed equally to this work.

## Abstract

Complex algorithms and enormous data sets require parallel execution of programs to attain results in a reasonable amount of time. Both aspects are combined in the domain of three-dimensional stencil operations, for example, computational fluid dynamics. This work contributes to the research on high-level parallel programming by discussing the generalizable implementation of a three-dimensional stencil skeleton that works in heterogeneous computing environments. Two exemplary programs, a gas simulation with the Lattice Boltzmann method, and a mean blur, are executed in a multi-node multi-graphics processing units (GPUs) environment, proving the runtime improvements in heterogeneous computing environments.

**Keywords:** Skeleton Programming, Three-Dimensional Stencil Operations

## 1 Introduction

The field of High Performance Computing (HPC) is growing as algorithms become more complex and more data is available. Evaluating massive datasets, therefore, requires writing efficient parallel programs. Most HPC environments have multiple nodes equipped with multiple central processing units (CPUs) and GPUs. Creating programs which combine multiple types of hardware requires knowledge of low-level frameworks such as Message Passing Interface (MPI) [1], OpenMP [2], and CUDA [3].

Writing a parallel program is error prone and tedious as e.g. out of memory errors, and invalid memory accesses are troublesome to identify even for skilled programmers. Moreover, choosing memory spaces, distributing data, and assigning task to threads are design decision which have a high impact on performance but require experience. Writing a program which combines multiple frameworks overstrains scientist.

Since experts in this field are hard to find, high-level frameworks are often used. Those frameworks commonly abstract from the distribution of data, provide portable code for different hardware architecture, are adjustable to distinct accelerators, and require less maintainance for the end-user. In 1989 COLE introduced algorithmic skeletons, enclosing reoccurring parallel and distributed computing patterns, such as Map and Reduce as one of the most common approaches to abstract from low-level details [4]. Multiple libraries [5, 6], general frameworks [7, 8], and domain-specific languages (DSLs)[9] utilize the concept.

The paper contributes to the ongoing work by focusing on a particularly arduous operation, namely three-dimensional stencil operations. Stencil operations calculate elements depending on other values inside the data structure and therefore require communication between the computational units used. Those operations are irreplaceable, for e.g., simulation of gas or computational fluid dynamics. Efficiently updating data in a generalizable way and dealing with 3D data structures are obstacles not solved in current high-level approaches.

This paper firstly elaborates on the related work, focusing on high-level approaches abstracting from problem specific details (Section 2). Section 3 outlines the library used (muesli), while Section 4 explains the additional implementation of the three-dimensional skeleton. The work is evaluated in Section 5 discussing our runtime experiments on multiple hardware set-ups. Lastly, Section 6 summarizes our work.

## 2 Related Work

Ongoing work discussing three-dimensional stencil operations is twofold. On the one hand, there exist generic high-level frameworks which have the advantage of parallelizing pre- and postprocessing steps, as many of them offer a variety of operations/skeletons. On the other hand, specialized frameworks already contain algorithms' implementations but are often inefficient. Most related regarding high-level skeleton programming, SkePU3 targets multi-node and multi-GPU environments for most skeletons in combination with StarPU. However, for stencil operations (MapOverlap), the exchange of data between the programs is missing for multi-node programs [7]. FastFlow added GPU support but focuses on communication skeletons and misses a comparable stencil operation [10][8].

Specialized libraries such as Palabos for the Lattice Boltzmann methods (LBMs) [11], or publications discussing a single method, e.g., the Helmholtz equation [12] focus rather on the algorithm and do not include accelerators as GPUs, which provide a significant speed-up.

This work extends the mentioned work as the presented stencil skeleton is generalizable for multiple methods and abstracts from the multiple layers of parallelism. This

is proven by showing the implementation of a LBM and a three-dimensional mean blur. Both programs run on multiple nodes and with multiple GPUs.

### 3 The *Muenster Skeleton Library Muesli*


```

1 class Sum : public Functor2<int, int, int>{
2   public: MSL_USERFUNC int operator() (int x, int y)
3     const {return x+y;}};
4 Sum sum;
5 auto product = [] (int i, int j) {return i*j;};
6 DA<int> a(3,2); // delivers: {2,2,2}
7 DA<int> b = a.mapIndex(sum); // delivers: {2,3,4}
8 a.zipInPlace(b,product); // delivers: {4,6,8}
9 int scalarproduct = a.fold(sum); // delivers: 18

```

### 4 Three-Dimensional Stencil Operations

Stencil operations are map operation which additionally require to read the surrounding elements of the datastructure. Figure 1 displays a two-dimensional stencil with the size one. The peculiarity regarding stencil operations on multiple nodes and accelerators is that each execution of the stencil operation requires to update elements which are shared between computational units. As communication of updated elements requires synchronization between the computational nodes it decreases the opportunity for executing task in parallel within the program. Muenster Skeleton Library (Muesli) abstracts from all communication between the computational nodes with a MapStencil Skeleton. The usage of the skeleton for the end-user is shown in Listing 1.



**Fig. 1:** Stencil Operation

#### 4.1 Using the MapStencil Skeleton

Firstly, a function which is executed on each element is defined (l. 1-12). Merely functors of type `DCMapStencilFunctor` are permitted to be used with the `mapStencil` skeleton. Therefore, the first argument of the functor has to be of type `PLCube` (*Padded-LocalCube*), and the following arguments have to be integers for indexing the datastructure. The class `PLCube` most importantly offers a getter taking three index arguments, relieving the end-user from index calculations (l.9). The presented functor calculates the sum of all elements with a radius of two and divides the sum with the number of total elements, therefore calculating a mean blur. This functor can be applied to a distributed cube by calling the `mapStencil` skeleton as a member function (l.19). The skeleton takes the functor as a template argument, and requires a distributed cube of the same dimension with the current data, the radius of the stencil and the neutral value for border elements.

Listing 1: Exemplary Functor for Stencil Skeleton

```

1  MSL_USERFUNC float update(const PLCube<float> &plCube,
2                          int x, int y, int z) {
3      float res = 0;
4      const int radius = 2;
5      const int elements = 48;
6      for (int mx = x - radius; mx <= x + radius; mx++) {
7          for (int my = y - radius; my <= y + radius; my++) {
8              for (int mz = z - radius; mz <= z + radius; mz++) {
9                  res += plCube(mx, my, mz);
10             }
11         }
12     }
13     return res/(elements);
14 }
15 ...
16 main () {
17     ...
18     int stencilradius = 2;
19     dcp1->mapStencil<update>(*dcp2, stencilradius, 0);
20     ...
21 }

```

## 4.2 Implementation of the MapStencil Skeleton

Adding the MapStencil skeleton to the existing distributed cube (DC) class requires to add two additional parameters: a vector of PLCubes and a supported stencil size. As previously mentioned the PLCubes class serves for the end-user to abstract from indexing of the datastructure. To make the access to the different memory spaces efficient each computational unit has a separate PLCubes storing merely the elements needed for the calculations of the assigned elements. This design choice makes the class flexible to be used for CPUs as well as for GPUs. It contains the following attributes to provide a light, minimal design:

- **int width, height, depth** the three dimensions of the datastructure,
- **int stencilSize** size of the Stencil required to calculate the overlapping elements,
- **int neutralValue** used when the index is outside the data structure,
- **T \* data, T\* topPadding, T\* bottomPadding** CPU or GPU pointer for current data,
- four integers to save global indexes for start and end of datastructure and stencil size.

Most importantly, the index operator is implemented, taking three integers as arguments returning suitable value. This is either the neutral value or the corresponding element from the CPU or GPU memory.

Assuming GPUs are used as accelerators, the skeleton updates the current datastructure in case the data is not up to date (Listing 2 l.6). Afterwards, it synchronizes the PLCubes inside one node, and the data between multiple nodes (l.7, l.9). For each GPU used the mapStencilKernelDC is called executing the functor on the appropriate part of the overall datastructure. In any other case (multiple nodes and cpu) it is only necessary to synchronize the nodes (l. 21), and thereafter call the functor with the corresponding arguments.

Listing 2: Implementation of MapStencil Skeleton

```

1  template<typename T>
2  template<msl::DCMapStencilFunctor<T> f>
3  void msl::DC<T>::mapStencil(msl::DC<T> &result, size_t stencilSize,
4                             T neutralValue) {
5      #ifdef __CUDACC__
6          this->updateDevice();
7          syncPLCubes(stencilSize, neutralValue);
8          msl::syncStreams();
9          syncPLCubesMPI(stencilSize);
10         for (int i = 0; i < this->ng; i++) {
11             cudaSetDevice(i);
12             dim3 dimBlock(Muesli::threads_per_block);
13             dim3 dimGrid((this->plans[i].size + dimBlock.x - 1) / dimBlock.x);
14             detail::mapStencilKernelDC<T, f><<<dimGrid, dimBlock, 0,
15                 Muesli::streams[i]>>>(result.plans[i].d_Data, this->plCubes[i],
16                 result.plans[i].size);
17         }
18         msl::syncStreams();
19         result.setCpuMemoryInSync(false);
20     #else
21         syncPLCubesMPI(stencilSize);
22     #ifdef _OPENMP
23     #pragma omp parallel for
24     #endif
25     for (int k = 0; k < this->nLocal; k++) {
26         int l = (k + this->firstIndex) / (ncol*nrow);
27         int j = ((k + this->firstIndex) - l*(ncol*nrow)) / ncol;
28         int i = (k + this->firstIndex) % ncol;
29         result.localPartition[k] = f(this->plCubes[0], i, j, l);
30     }
31     #endif
32 }

```

### 4.3 Example Application for Three-Dimensional Stencil Operations

For the evaluation of our implementation two examples were used: a LBM implementation and a mean blur. The exemplary user function of the mean blur was already shown in Listing 1. However, the implementation of a LBM underlines the applicability for real application contexts.

LBM is used for fluid simulations. It distinguishes between the collision and the streaming step which alternate in continuous simulations [13, p.61ff]. In the streaming step particles move from one cell to another. In the collision step the fluid flow caused by the colliding particles is calculated. The distribution function  $f_i(x, t)$  calculates for a cell  $x$  and a timestamp  $t$  how many particles move in the next step to the neighbour  $i$ . Zero is the cell itself.

$$f_i(x + c_i \Delta t, t + \Delta t) := f_i^*(x, t) \quad (1)$$

For the collision steps the Bhatnagar-Gross-Krook-operator is used.  $f_i^*$  defines the distribution after the collision of the particles,  $\Delta t$  the time period to be simulated and  $\tau$  a constant defining the convergence of the simulation. Thus  $\tau$  influences the viscosity of the gases.

$$f_i^*(x, t) := f_i(x, t) - \frac{\Delta t}{\tau} (f_i(x, t) - f_i^{\text{eq}}(x, t)). \quad (2)$$

The equilibrium state is calculated by

$$f_i^{\text{eq}}(x, t) := w_i \rho \left( 1 + \frac{u \cdot c_i}{c_s^2} + \frac{u \cdot c_i}{2c_s^4} + \frac{u \cdot u}{2c_s^2} \right), \quad (3)$$

where  $w_i$  are the weights of the chosen grid and  $c_i$  is the position of the neighbour cells relative to the main cell. The constant number  $c_s$  is the sound velocity of the model. The mass density  $\rho$  and the puls density  $u$  are defined by

$$\rho(x, t) = \sum_i f_i(x, t), \quad \rho u(x, t) := \sum_i c_i f_i(x, t). \quad (4)$$

For the implementation of the LBM a D3Q19-Grid was used, D being the number of dimensions and Q the number of neighbours. Both steps (collision and streaming) are combined in one `mapStencil` call. Noteworthy, the implementation has to take into consideration that single cells can be marked as blocked, simulating objects which are barriers for the flow of air or as distributing constant velocity. Therefore, special cells are marked with `Not a Number` values (Listing 4 l. 5-7). To simulate this behaviour without requiring additional storage, the handling of the floating numbers is extended. According to the IEEE-754 Standard each float has a maximal exponent with a mantissa which is not equal to zero. To identify special cells the most significant bit of the mantissa of `f0` is set, so that the number is definitely understood as a `NaN`. The remaining bits of the mantissa can then be used freely to store other data.

In the code, bit masks and a struct with bit- fields are defined in the code to access this information as easily as possible (Listing 3).

Listing 3: Handling of Barriers and Streaming Cells

```

1  const int FLAG_OBSTACLE = 1 << 0;
2  const int FLAG_KEEP_VELOCITY = 1 << 1;
3  typedef struct {
4      unsigned int mantissa : 23;
5      unsigned int exponent : 8;
6      unsigned int sign : 1;
7  } floatparts ;

```

The data stored for each cell is an `array<float, Q>`. `Q` is a constant number for the neighbour cells and the cell itself (19). This type is appreciated in the following listing with `cell_t`. Moreover, it is abstracted from the three-dimensional vector operations (1. 28,29,31,34). The user function starts by transforming the current value of the cell into the single float parts (1.4). In case we have a cell which distributes gas (`FLAG_KEEP_VELOCITY`), the cell remains without changes (1.5-7). For all neighbour cells the current amount of particles is read (1.10-12). In the collision step all cells which are obstacles reverse the flow of air (1.16-23). All other cells calculate the particles streaming from the next cells.

Listing 4: Implementation of a Exemplary LBM User Function

```

1  MSL_USERFUNC cell_t update(const PLCube<cell_t> &plCube, int x,
2                               int y, int z) {
3      cell_t cell = plCube(x, y, z);
4      auto* parts = (floatparts*) &cell[0];
5      if (parts->exponent == 255 && parts->mantissa
6          & FLAG_KEEP_VELOCITY) {
7          return cell;
8      }
9      // Streaming.
10     for (int i = 1; i < Q; i++) {
11         cell[i] = plCube(x + (int) offsets[i].x,
12                         y + (int) offsets[i].y, z + (int) offsets[i].z)[i];
13     }
14
15     // Collision.
16     if (parts->exponent == 255 && parts->mantissa & FLAG_OBSTACLE) {
17         if (parts->mantissa & FLAG_OBSTACLE) {
18             cell_t cell2 = cell;
19             for (size_t i = 1; i < Q; i++) {
20                 cell[i] = cell2[opposite[i]];
21             }
22         }
23         return cell;

```

Number Nodes	GPU-type	Per Node		
		GPUs	CPU-type	CPUs
5	GeForce RTX 2080 Ti	8	Zen3(EPYC 7513) 24 cores	1
2	Nvidia A100 SXM 80GB	8	Zen3(EPYC 7513) 24 cores	1
12	-	-	Zen2 (EPYC 7742) 24 cores	128

**Table 1:** Overview of used Hardware

```

24  }
25  float p = 0;
26  vec3f vp {0, 0, 0};
27  for (size_t i = 0; i < Q; i++) {
28      p += cell[i];
29      vp += offsets[i] * cellwidth * cell[i];
30  }
31  vec3f v = p == 0 ? vp : vp * (1 / p);
32
33  for (size_t i = 0; i < Q; i++) {
34      cell[i] = cell[i] + deltaT / tau * (freq(i, p, v) - cell[i]);
35  }
36  return cell;
37 }

```

## 5 Evaluation


For the experiments, the HPC machine Palma II<sup>1</sup>. We used the GeForce RTX 2080 Ti GPUs partition equipped with 5 nodes each with 8 GPUs and a Zen3 (EPYC 7513) CPU. Complementary we used the gpubgx partition equipped with 2 Nodes each with 8 A100 SXM GPUs. For testing CPU-parallelization we used the smp Partition equipped with 12 nodes with a Zen2 (EPYC 7742) CPUs with 64 cores. To provide generalizable results, the mean runtime of 20 execution was used. For running the sequential version on the HPC, a single Broadwell (E5-2683 v4) CPU was used.

## 6 Conclusion


Equations in  $\text{\LaTeX}$  can either be inline or on-a-line by itself (“display equations”). For inline equations use the  $\$...\$$  commands. E.g.: The equation  $H\psi = E\psi$  is written via the command  $\$H \backslash\psi = E \backslash\psi\$$ .

<sup>1</sup><https://confluence.uni-muenster.de/display/HPC/GPU+Nodes>





**Fig. 2:** This is a widefig. This is an example of long caption this is an example of long caption this is an example of long caption this is an example of long caption



**Fig. 3:** This is a widefig. This is an example of long caption this is an example of long caption this is an example of long caption this is an example of long caption


For display equations (with auto generated equation numbers) one can use the equation or align environments:

$$\|\tilde{X}(k)\|^2 \leq \frac{\sum_{i=1}^p \|\tilde{Y}_i(k)\|^2 + \sum_{j=1}^q \|\tilde{Z}_j(k)\|^2}{p+q}. \quad (5)$$

where,

$$\begin{aligned} D_\mu &= \partial_\mu - ig \frac{\lambda^a}{2} A_\mu^a \\ F_{\mu\nu}^a &= \partial_\mu A_\nu^a - \partial_\nu A_\mu^a + gf^{abc} A_\mu^b A_\nu^c \end{aligned} \quad (6)$$

Notice the use of `\nonumber` in the align environment at the end of each line, except the last, so as not to produce equation numbers on lines where no equation numbers are required. The `\label{}` command should only be used at the last line of an align



**Fig. 4:** This is a widefig. This is an example of long caption this is an example of long caption this is an example of long caption this is an example of long caption

environment where `\nonumber` is not used.

$$Y_{\infty} = \left( \frac{m}{\text{GeV}} \right)^{-3} \left[ 1 + \frac{3 \ln(m/\text{GeV})}{15} + \frac{\ln(c_2/5)}{15} \right] \quad (7)$$

The class file also supports the use of `\mathbb{}`, `\mathscr{}` and `\mathcal{}` commands. As such `\mathbb{R}`, `\mathscr{R}` and `\mathcal{R}` produces  $\mathbb{R}$ ,  $\mathscr{R}$  and  $\mathcal{R}$  respectively (refer Subsubsection ??).

Tables can be inserted via the normal table and tabular environment. To put footnotes inside tables you should use `\footnotetext[]{\dots}` tag. The footnote appears just below the table itself (refer Tables 2 and 3). For the corresponding footnote mark use `\footnotemark[...]`

The input format for the above table is as follows:

```
\begin{table}[<placement-specifier>]
\caption{<table-caption>}\label{<table-label>}%
\begin{tabular}{@{}llll@{}}
\toprule
Column 1 & Column 2 & Column 3 & Column 4\\
\midrule
row 1 & data 1 & data 2 & data 3 \\
row 2 & data 4 & data 5\footnotemark[1] & data 6 \\
row 3 & data 7 & data 8 & data 9\footnotemark[2]\end{pre>
```

**Table 2:** Caption text

Column 1	Column 2	Column 3	Column 4
row 1	data 1	data 2	data 3
row 2	data 4	data 5 <sup>1</sup>	data 6
row 3	data 7	data 8	data 9 <sup>2</sup>

Source: This is an example of table footnote. This is an example of table footnote.

<sup>1</sup>Example for a first table footnote. This is an example of table footnote.

<sup>2</sup>Example for a second table footnote. This is an example of table footnote.

```
\botrule
\end{tabular}
\footnotetext{Source: This is an example of table footnote.
This is an example of table footnote.}
\footnotetext[1]{Example for a first table footnote.
This is an example of table footnote.}
\footnotetext[2]{Example for a second table footnote.
This is an example of table footnote.}
\end{table}
```

**Table 3:** Example of a lengthy table which is set to full textwidth

Project	Element 1 <sup>1</sup>			Element 2 <sup>2</sup>		
	Energy	$\sigma_{calc}$	$\sigma_{expt}$	Energy	$\sigma_{calc}$	$\sigma_{expt}$
Element 3	990 A	1168	$1547 \pm 12$	780 A	1166	$1239 \pm 100$
Element 4	500 A	961	$922 \pm 10$	900 A	1268	$1092 \pm 40$

Note: This is an example of table footnote. This is an example of table footnote this is an example of table footnote this is an example of table footnote this is an example of table footnote.

<sup>1</sup>Example for a first table footnote.

<sup>2</sup>Example for a second table footnote.

In case of double column layout, tables which do not fit in single column width should be set to full text width. For this, you need to use `\begin{table*} ... \end{table*}` instead of `\begin{table} ... \end{table}` environment. Lengthy tables which do not fit in textwidth should be set as rotated table. For this, you need to use `\begin{sidewaystable} ... \end{sidewaystable}` instead of `\begin{table*} ... \end{table*}` environment. This environment puts tables rotated to single column width. For tables rotated to double column width, use `\begin{sidewaystable*} ... \end{sidewaystable*}`.

As per the  $\text{\LaTeX}$  standards you need to use eps images for  $\text{\LaTeX}$  compilation and pdf/jpg/png images for PDF $\text{\LaTeX}$  compilation. This is one of the major difference between  $\text{\LaTeX}$  and PDF $\text{\LaTeX}$ . Each image should be from a single input .eps/vector image file. Avoid using subfigures. The command for inserting images for  $\text{\LaTeX}$  and PDF $\text{\LaTeX}$  can be generalized. The package used to insert images in  $\text{\LaTeX}$ /PDF $\text{\LaTeX}$  is the graphicx package. Figures can be inserted via the normal figure environment as shown in the below example:

```
\begin{figure}[<placement-specifier>]
\centering
\includegraphics{<eps-file>}
\caption{<figure-caption>}\label{<figure-label>}
\end{figure}
```

In case of double column layout, the above format puts figure captions/images to single column width. To get spanned images, we need to provide `\begin{figure*} ... \end{figure*}`.

For sample purpose, we have included the width of images in the optional argument of `\includegraphics` tag. Please ignore this.

## 7 Algorithms, Program codes and Listings

Packages `algorithm`, `algorithmicx` and `algpseudocode` are used for setting algorithms in  $\text{\LaTeX}$  using the format:

```
\begin{algorithm}
\caption{<alg-caption>}\label{<alg-label>}
\begin{algorithmic}[1]
. . .
\end{algorithmic}
\end{algorithm}
```

You may refer above listed package documentations for more details before setting `algorithm` environment. For program codes, the “verbatim” package is required and the command to be used is `\begin{verbatim} ... \end{verbatim}`.

Similarly, for listings, use the `listings` package. `\begin{lstlisting} ... \end{lstlisting}` is used to set environments similar to `verbatim` environment. Refer to the `lstlisting` package documentation for more details.

A fast exponentiation procedure:

**Table 4:** Tables which are too long to fit, should be written using the “sidewaystable” environment as shown here

Projectile	Element 1 <sup>1</sup>				Element <sup>2</sup>	
	Energy	$\sigma_{calc}$	$\sigma_{expt}$	Energy	$\sigma_{calc}$	$\sigma_{expt}$
Element 3	990 A	1168	1547 $\pm$ 12	780 A	1166	1239 $\pm$ 100
Element 4	500 A	961	922 $\pm$ 10	900 A	1268	1092 $\pm$ 40
Element 5	990 A	1168	1547 $\pm$ 12	780 A	1166	1239 $\pm$ 100
Element 6	500 A	961	922 $\pm$ 10	900 A	1268	1092 $\pm$ 40

Note: This is an example of table footnote this is an example of table footnote this is an example of table footnote this is an example of table footnote  
this is an example of table footnote.

<sup>1</sup>This is an example of table footnote.

```

1  begin
2    for  $i := 1$  to 10 step 1 do
3       $\text{expt}(2, i)$ ;
4       $\text{newline}()$  od
5  where
6  proc  $\text{expt}(x, n) \equiv$ 
7     $z := 1$ ;
8    do if  $n = 0$  then  $\text{exit fi}$ ;
9      do if  $\text{odd}(n)$  then  $\text{exit fi}$ ;
10      $\text{comment: This is a comment statement;}$ 
11      $n := n/2$ ;  $x := x * x$  od;
12     {  $n > 0$  };
13      $n := n - 1$ ;  $z := z * x$  od;
14    $\text{print}(z)$ .
15 end

```

Comments will be set flush to the right margin

---

**Algorithm 1** Calculate  $y = x^n$

---

**Require:**  $n \geq 0 \vee x \neq 0$

**Ensure:**  $y = x^n$

```

1:  $y \leftarrow 1$ 
2: if  $n < 0$  then
3:    $X \leftarrow 1/x$ 
4:    $N \leftarrow -n$ 
5: else
6:    $X \leftarrow x$ 
7:    $N \leftarrow n$ 
8: end if
9: while  $N \neq 0$  do
10:  if  $N$  is even then
11:     $X \leftarrow X \times X$ 
12:     $N \leftarrow N/2$ 
13:  else [ $N$  is odd]
14:     $y \leftarrow y \times X$ 
15:     $N \leftarrow N - 1$ 
16:  end if
17: end while

```

---

```

1  for i:=maxint to 0 do
2  begin
3  { do nothing }
4  end;
5  Write('Case insensitive ');
6  Write('Pascal keywords. ');

```

## 8 Cross referencing

Environments such as figure, table, equation and align can have a label declared via the `\label{#label}` command. For figures and table environments use the `\label{}` command inside or just below the `\caption{}` command. You can then use the `\ref{#label}` command to cross-reference them. As an example, consider the label declared for Figure 4 which is `\label{fig1}`. To cross-reference it, use the command `\ref{fig1}`, for which it comes up as “Figure 4”.

To reference line numbers in an algorithm, consider the label declared for the line number 2 of Algorithm 1 is `\label{algn2}`. To cross-reference it, use the command `\ref{algn2}` for which it comes up as line 2 of Algorithm 1.

### 8.1 Details on reference citations

Standard  $\text{\LaTeX}$  permits only numerical citations. To support both numerical and author-year citations this template uses `natbib`  $\text{\LaTeX}$  package. For style guidance please refer to the template user manual.

Here is an example for `\cite{...}`: [1]. Another example for `\citep{...}`: [1]. For author-year citation mode, `\cite{...}` prints Jones et al. (1990) and `\citep{...}` prints (Jones et al., 1990).

All cited bib entries are printed at the end of this article: [1], [1], [1].

## 9 Examples for theorem like environments

For theorem like environments, we require `amsthm` package. There are three types of predefined theorem styles exists—`thmstyleone`, `thmstyletwo` and `thmstylethree`

<code>thmstyleone</code>	Numbered, theorem head in bold font and theorem text in italic style
<code>thmstyletwo</code>	Numbered, theorem head in roman font and theorem text in italic style
<code>thmstylethree</code>	Numbered, theorem head in bold font and theorem text in roman style

For mathematics journals, theorem styles can be included as shown in the following examples:

**Theorem 1** (Theorem subhead). *Example theorem text. Example theorem text. Example theorem text. Example theorem text. Example theorem text. Example theorem text. Example theorem text. Example theorem text. Example theorem text. Example theorem text.*

Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text.

**Proposition 2.** *Example proposition text. Example proposition text. Example proposition text. Example proposition text. Example proposition text. Example proposition text. Example proposition text. Example proposition text. Example proposition text.*

Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text.

**Example 1.** *Phasellus adipiscing semper elit. Proin fermentum massa ac quam. Sed diam turpis, molestie vitae, placerat a, molestie nec, leo. Maecenas lacinia. Nam ipsum ligula, eleifend at, accumsan nec, suscipit a, ipsum. Morbi blandit ligula feugiat magna. Nunc eleifend consequat lorem.*

Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text.

**Remark 1.** *Phasellus adipiscing semper elit. Proin fermentum massa ac quam. Sed diam turpis, molestie vitae, placerat a, molestie nec, leo. Maecenas lacinia. Nam ipsum ligula, eleifend at, accumsan nec, suscipit a, ipsum. Morbi blandit ligula feugiat magna. Nunc eleifend consequat lorem.*

Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text.

**Definition 1** (Definition sub head). *Example definition text. Example definition text. Example definition text. Example definition text. Example definition text. Example definition text.*

Additionally a predefined “proof” environment is available: `\begin{proof}` ... `\end{proof}`. This prints a “Proof” head in italic font style and the “body text” in roman font style with an open square at the end of each proof environment.

*Proof.* Example for proof text. Example for proof text. Example for proof text. Example for proof text. Example for proof text. Example for proof text. Example for proof text. Example for proof text. □

Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text.

*Proof of Theorem 1.* Example for proof text. Example for proof text. Example for proof text. Example for proof text. Example for proof text. Example for proof text. Example for proof text. Example for proof text. □

For a quote environment, use `\begin{quote}` ... `\end{quote}`

Quoted text example. Aliquam porttitor quam a lacus. Praesent vel arcu ut tortor cursus volutpat. In vitae pede quis diam bibendum placerat. Fusce elementum convallis neque.



Sed dolor orci, scelerisque ac, dapibus nec, ultricies ut, mi. Duis nec dui quis leo sagittis commodo.

Sample body text. Sample body text. Sample body text. Sample body text. Sample body text (refer Figure 4). Sample body text. Sample body text. Sample body text (refer Table 4).

## 10 Methods

Topical subheadings are allowed. Authors must ensure that their Methods section includes adequate experimental and characterization data necessary for others in the field to reproduce their work. Authors are encouraged to include RIIDs where appropriate.

**Ethical approval declarations** (only required where applicable) Any article reporting experiment/s carried out on (i) live vertebrate (or higher invertebrates), (ii) humans or (iii) human samples must include an unambiguous statement within the methods section that meets the following requirements:

1. Approval: a statement which confirms that all experimental protocols were approved by a named institutional and/or licensing committee. Please identify the approving body in the methods section
2. Accordance: a statement explicitly saying that the methods were carried out in accordance with the relevant guidelines and regulations
3. Informed consent (for experiments involving humans or human tissue samples): include a statement confirming that informed consent was obtained from all participants and/or their legal guardian/s

If your manuscript includes potentially identifying patient/participant information, or if it describes human transplantation research, or if it reports results of a clinical trial then additional information will be required. Please visit (<https://www.nature.com/nature-research/editorial-policies>) for Nature Portfolio journals, (<https://www.springer.com/gp/authors-editors/journal-author/journal-author-helpdesk/publishing-ethics/14214>) for Springer Nature journals, or (<https://www.biomedcentral.com/getpublished/editorial-policies#ethics+and+consent>) for BMC.

## 11 Discussion

Discussions should be brief and focused. In some disciplines use of Discussion or ‘Conclusion’ is interchangeable. It is not mandatory to use both. Some journals prefer a section ‘Results and Discussion’ followed by a section ‘Conclusion’. Please refer to Journal-level guidance for any specific requirements.

## 12 Conclusion

Conclusions may be used to restate your hypothesis or research question, restate your major findings, explain the relevance and the added value of your work, highlight any limitations of your study, describe future directions for research and recommendations.

In some disciplines use of Discussion or 'Conclusion' is interchangeable. It is not mandatory to use both. Please refer to Journal-level guidance for any specific requirements.

**Supplementary information.** If your article has accompanying supplementary file/s please state so here.

Authors reporting data from electrophoretic gels and blots should supply the full unprocessed scans for key as part of their Supplementary information. This may be requested by the editorial team/s if it is missing.

Please refer to Journal-level guidance for any specific requirements.

**Acknowledgments.** Acknowledgments are not compulsory. Where included they should be brief. Grant or contribution numbers may be acknowledged.

Please refer to Journal-level guidance for any specific requirements.

## Declarations

Some journals require declarations to be submitted in a standardised format. Please check the Instructions for Authors of the journal to which you are submitting to see if you need to complete this section. If yes, your manuscript must contain the following sections under the heading 'Declarations':

- Funding
- Conflict of interest/Competing interests (check journal-specific guidelines for which heading to use)
- Ethics approval
- Consent to participate
- Consent for publication
- Availability of data and materials
- Code availability
- Authors' contributions

If any of the sections are not relevant to your manuscript, please include the heading and write 'Not applicable' for that section.

Editorial Policies for:

Springer journals and proceedings: <https://www.springer.com/gp/editorial-policies>

Nature Portfolio journals:  
<https://www.nature.com/nature-research/editorial-policies>

*Scientific Reports*: <https://www.nature.com/srep/journal-policies/editorial-policies>

BMC journals: <https://www.biomedcentral.com/getpublished/editorial-policies>

## Appendix A Section title of first appendix

An appendix contains supplementary information that is not an essential part of the text itself but which may be helpful in providing a more comprehensive understanding of the research problem or it is information that is too cumbersome to be included in the body of the paper.

## References

- [1] Forum, M.: MPI Standard. <https://www.mpi-forum.org/docs/>. Accessed: 10.05.2021 (2021)
- [2] OpenMP: OpenMP The OpenMP API specification for parallel programming. <https://www.openmp.org/>. Accessed: 10.05.2021 (2021)
- [3] Corporation, N.: CUDA. <https://developer.nvidia.com/cuda-zone>. Accessed: 10.05.2021 (2021)
- [4] Cole, M.I.: Algorithmic Skeletons: Structured Management of Parallel Computation. Pitman London, ??? (1989)
- [5] Ernsting, S., Kuchen, H.: Data parallel algorithmic skeletons with accelerator support. *International Journal of Parallel Programming* **45**(2), 283–299 (2017)
- [6] Benoit, A., Cole, M., Gilmore, S., Hillston, J.: Flexible skeletal programming with eskel. In: *European Conference on Parallel Processing*, pp. 761–770 (2005). Springer
- [7] Ernstsson, A., Ahlqvist, J., Zouzoula, S., Kessler, C.: Skepu 3: Portable high-level programming of heterogeneous systems and hpc clusters. *International Journal of Parallel Programming* **49**(6), 846–866 (2021)
- [8] Aldinucci, M., Danelutto, M., Kilpatrick, P., Torquati, M.: Fastflow: high-level and efficient streaming on multi-core. *Programming multi-core and many-core computing systems, parallel and distributed computing* (2017)
- [9] Wrede, F., Rieger, C., Kuchen, H.: Generation of high-performance code based on a domain-specific language for algorithmic skeletons. *The Journal of Supercomputing* **76**(7), 5098–5116 (2020)
- [10] Goli, M., González-Vélez, H.: Heterogeneous algorithmic skeletons for fast flow with seamless coordination over hybrid architectures. In: *2013 21st Euromicro International Conference on Parallel, Distributed, and Network-Based Processing*, pp. 148–156 (2013). <https://doi.org/10.1109/PDP.2013.29>
- [11] Latt, J., Malaspinas, O., Kontaxakis, D., Parmigiani, A., Lagrava, D., Brogi, F.,

- Belgacem, M.B., Thorimbert, Y., Leclaire, S., Li, S., Marson, F., Lemus, J., Kotsalos, C., Conradin, R., Coreixas, C., Petkantchin, R., Raynaud, F., Beny, J., Chopard, B.: Palabos: Parallel lattice boltzmann solver. *Computers and Mathematics with Applications* **81**, 334–350 (2021) <https://doi.org/10.1016/j.camwa.2020.03.022>. Development and Application of Open-source Software for Problems with Numerical PDEs
- [12] Gonzales, R., Gryazin, Y., Lee, Y.T.: Parallel fft algorithms for high-order approximations on three-dimensional compact stencils. *Parallel Computing* **103**, 102757 (2021) <https://doi.org/10.1016/j.parco.2021.102757>
- [13] Krüger, T., Kusumaatmaja, H., Kuzmin, A., Shardt, O., Silva, G., Viggen, E.M.: *The Lattice Boltzmann Method: Principles and Practice*. Graduate Texts in Physics. Springer, Cham (2016)